

CS559: Computer Systems Design — Lab Report

Student Name: Pankaj Kashyap

Assigned Host: 10.1.75.51

Course: CS559 Computer Systems Design

Roll Number: 230101053

Submitted LB URL: <http://10.1.75.51:3000>

Repository: github.com/Pankajkashyap1/Groupchat-wavetalk

1. SYSTEM DEPLOYMENT & ALLOTTED PORTS

The complete distributed messaging cluster is deployed across the 4 allotted laboratory environments on host 10.1.75.51. Clients, benchmark tools, and the official evaluation load generator interact exclusively with the **Sys1 Load Balancer** on port 3000.

System	Role	Host / IP	SSH Port	App Port	Description
Sys1	Load Balancer	10.1.75.51	2257	3000	Dynamic reverse proxy, exposes /message & /feed, health check
Sys2	Backend 1	10.1.75.51	2258	3001	WaveTalk chat service, shared DB persistence, /api/metrics
Sys3	Backend 2	10.1.75.51	2259	3002	WaveTalk chat service, shared DB persistence, /api/metrics
Sys4	Backend 3	10.1.75.51	2260	3003	WaveTalk chat service, shared DB persistence, /api/metrics

2. PERFORMANCE-BASED DYNAMIC LOAD BALANCING

Instead of fixed round-robin, the Load Balancer evaluates real-time performance telemetry collected from all backends every 2.5 seconds. The load score is computed continuously as:

```
Load Score = (ActiveConnections × 0.40) + (CPULoad% × 0.40) + (AvgResponseTime / 20 × 0.20) Routing Rule: If CurrentBackend.Score > 65 (Threshold), switch to min(HealthyBackends.Score)
```

Dynamic Switching: The current backend remains active while its score is below the threshold (65). When load rises above 65, new requests switch immediately to the least-loaded healthy backend.

Health Monitoring: Any backend failing 3 consecutive health probes is evicted from the active pool and traffic fails over automatically. Recovered backends rejoin the pool with zero downtime.

3. SHARED DATABASE & DUPLICATE PREVENTION

- Shared Persistence:** Backends share /home/student/shared_db/chat_database.json.
- Idempotency & No Duplicates:** Every message has a unique message ID (client UUID or deterministic ID). An in-memory Set index detects duplicates in O(1) time. Redundant submissions are ignored and return status: "duplicate" without double insertion.
- Atomic File Writes:** State is saved via .tmp.[timestamp] atomic rename, preventing file corruption under concurrent writes.

4. REQUIRED API ROUTES

Route	Method	Input	Output
/message	POST / GET	client-name, msg	{ "id": "...", "status": "ok" "duplicate", "sender": "...", "msg": "...", "timestamp": ... }
/feed	GET	None	[{ "id": "...", "sender": "...", "msg": "...", "timestamp": ... }] (All messages)

5. OPTIMAL SWITCHING THRESHOLD ANALYSIS

A threshold parameter sweep was conducted using the automated load generator with 100 concurrent workers to identify the optimal switching point:

Threshold	Throughput (RPS)	Avg Latency	P95 Latency	Switches/Min	Observation
40	172.4 req/s	64.2 ms	118.5 ms	48	Too sensitive; excessive switching churn
50	198.1 req/s	55.8 ms	99.2 ms	28	Acceptable, but premature backend shifts
65 (Optimal)	224.8 req/s	46.3 ms	82.4 ms	14	Optimal: peak throughput & low tail latency
80	188.5 req/s	62.7 ms	124.1 ms	6	Sluggish; backend queue buildup prior to switch
90	154.2 req/s	84.1 ms	168.9 ms	2	Overloaded; severe queuing delay before shift

6. RESPONSE TIME & 4-SYSTEM UTILIZATION PLOTS




Figure 1: Average & P95 Latency across Concurrency Tiers




Figure 2: CPU Utilization (%) across Sys1 (LB), Sys2, Sys3, and Sys4

7. REFERENCES & AI ASSISTANCE DISCLOSURE

Artificial Intelligence Reference:
This project utilized **Google Gemini AI (via Google DeepMind Antigravity)** as an AI pair-programming assistant for:

- Designing the adaptive multi-metric scoring formula (combining CPU, active connections, and latency).
- Formulating the threshold optimization methodology and comparative switching stability analysis.
- Developing the variable load generator benchmarking engine and generating system utilization SVG plots.
- Drafting and formatting this concise submission documentation and architecture verification checklist.

Conclusion: All required features (dynamic threshold switching, health eviction, shared database with deduplication, and exact /message & /feed routes) are implemented and verified. The application maintains all full secure chat capabilities from the previous assignment.