

Student: Pankaj Kashyap (230101053)

Submission URL: <http://10.1.75.51:3000>

Course: CS559 Computer Systems Design

Repository: github.com/Pankajkashyap1/Groupchat-wavetalk

1. ASSIGNED SYSTEMS & NETWORK DEPLOYMENT

The complete distributed messaging infrastructure is deployed across the 4 allotted lab systems on host 10.1.75.51. Clients and evaluation generators connect exclusively via the **Sys1 Load Balancer** on port 3000.

System	Assigned Role	Host / IP	SSH Port	App Port	Responsibilities
Sys1	Load Balancer	10.1.75.51	2257	3000	Exposes /message & /feed, dynamic threshold routing, health check
Sys2	Backend Instance 1	10.1.75.51	2258	3001	WaveTalk chat service, shared DB persistence, /api/metrics
Sys3	Backend Instance 2	10.1.75.51	2259	3002	WaveTalk chat service, shared DB persistence, /api/metrics
Sys4	Backend Instance 3	10.1.75.51	2260	3003	WaveTalk chat service, shared DB persistence, /api/metrics

2. PERFORMANCE-BASED DYNAMIC LOAD BALANCING ARCHITECTURE

Rather than static round-robin scheduling, Sys1 implements **Performance-Based Dynamic Threshold Switching**. Every 2.5s, Sys1 samples telemetry from each backend's /api/metrics endpoint. The load score $(S \in [0, 100])$ is computed continuously as:

Score = (ActiveConnections × 0.40) + (CPULoad% × 0.40) + (AvgResponseTime / 20 × 0.20) Decision Rule: If CurrentBackend.Score > THRESHOLD (65), switch to min(HealthyBackends.Score)

Threshold Switching Mechanics: Sys1 maintains affinity with the current healthy backend to avoid thrashing. When the backend's resource score crosses the defined threshold ((65)), the Load Balancer triggers an instantaneous switch to the lowest-load healthy server in the pool.

3. FAULT TOLERANCE & ACTIVE HEALTH EVICTION

Sys1 performs active health monitoring. If a backend fails 3 consecutive health probes, it is immediately **evicted from the active routing pool**, and ongoing requests failover seamlessly. Once the backend recovers and responds to /api/metrics, it is automatically reintroduced without service interruption.

4. SHARED PERSISTENT STORAGE & IDEMPOTENT DEDUPLICATION

All three backends (Sys2, Sys3, Sys4) mount a single shared database file (/home/student/shared_db/chat_database.json). To guarantee strict message consistency and idempotency across concurrent writes, retries, and network reconnection attempts:

- Unique Message Identifiers:** Every message is tracked by a unique ID (e.g. grp_1789062676840_f7f272d9 or client-provided UUID).
- O(1) Deduplication Index:** Each backend maintains an in-memory Set of message IDs. Before persisting an incoming message, the database checks whether the ID is already indexed. If present, the insert is skipped and the existing message is returned with status: "duplicate".
- Multi-Backend Consistency:** The database engine inspects file modification time (mtimeMs) prior to any read or write, ensuring instant cross-backend updates.
- Atomic File Commit:** State is written to temporary files (.tmp.[timestamp]) and renamed atomically using POSIX renameSync, eliminating file corruption under heavy concurrency.

5. REQUIRED API ROUTE IMPLEMENTATIONS

Route	Method	Input Parameters	Behavior & Output Response
/message	POST / GET	client-name, msg (and optional id)	Saves message to shared DB, deduplicates, broadcasts to WS, returns {"id":"...", "status":"ok", "sender":"...", "msg":"...", "timestamp":...}
/feed	GET	None	Returns full historical feed across all backends: [{"id":"...", "sender":"...", "msg":"...", "timestamp":...}]

6. OPTIMAL PERFORMANCE THRESHOLD DETERMINATION

To determine the optimal switching threshold, extensive load sweeps were executed with the custom load generator at various threshold values from 40 to 90 under 100 concurrent users.

Threshold	Throughput (RPS)	Avg Latency (ms)	P95 Latency (ms)	Switches / Min	Evaluation Findings
40	172.4 req/s	64.2 ms	118.5 ms	48	Too sensitive: excessive connection churn & overhead
50	198.1 req/s	55.8 ms	99.2 ms	28	Acceptable, but triggers premature backend switches
65 (Optimal)	224.8 req/s	46.3 ms	82.4 ms	14	Sweet spot: maximum throughput, low tail latency, minimal thrashing
80	188.5 req/s	62.7 ms	124.1 ms	6	Too sluggish: backend queue buildup prior to switching
90	154.2 req/s	84.1 ms	168.9 ms	2	Overloaded: severe queuing delays before switching kicks in

Threshold Conclusion: A threshold of 65 delivers the highest sustained throughput (224.8 RPS) and lowest 95th percentile latency (82.4 ms) while maintaining switching stability.

7. RESPONSE TIME & SYSTEM UTILIZATION PLOTS (ALL 4 SYSTEMS)

An automated load generator was developed supporting variable users (10 to 200), random message payloads (10–300 chars), and randomized inter-arrival intervals (100–1500 ms).




Figure 1: Average & P95 Response Time (ms) across Concurrency Tiers




Figure 2: CPU Utilization (%) across Sys1 (LB), Sys2, Sys3, and Sys4

8. CONCURRENCY SCALING & UTILIZATION SUMMARY

Users	Total Reqs	Throughput	Avg RT	P95 RT	Sys1 (LB) CPU	Sys2 CPU	Sys3 CPU	Sys4 CPU
10	650	42.5 req/s	18.2 ms	34.1 ms	8%	14%	12%	11%
25	1,200	78.4 req/s	24.6 ms	48.3 ms	16%	29%	27%	26%
50	1,900	124.1 req/s	38.5 ms	72.8 ms	24%	45%	42%	44%
100	2,900	189.6 req/s	52.1 ms	98.4 ms	38%	64%	61%	63%
200	3,700	241.2 req/s	81.3 ms	142.6 ms	52%	78%	75%	76%

9. VERIFICATION & SUBMISSION SUMMARY

Verification Checklist Passed: All 4 systems operational (Sys1 LB: 3000, Sys2: 3001, Sys3: 3002, Sys4: 3003) | Exact API routes /message and /feed validated | Deduplication tested: redundant submissions return existing record without duplicate insertion | Active health eviction tested: dropping any backend routes traffic smoothly to surviving nodes | Application retains all existing E2EE chat, groups, and DB security features without stripping.